

```
In [17]: import pandas as pd
        import numpy as np
        import matplotlib.pyplot as plt
        import seaborn as sns

        from sklearn.model_selection import train_test_split
        from sklearn.ensemble import RandomForestClassifier
        from sklearn.linear_model import SGDClassifier
        from sklearn.svm import SVC

        from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
In [18]: df = pd.read_csv("WineQt.csv")
          df.head()
```

|   | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides | free sulfur dioxide | total sulfur dioxide | density | pH   | sulphates | alcohol |
|---|---------------|------------------|-------------|----------------|-----------|---------------------|----------------------|---------|------|-----------|---------|
| 0 | 7.4           | 0.70             | 0.00        | 1.9            | 0.076     | 11.0                | 34.0                 | 0.9978  | 3.51 | 0.56      | 10.5    |
| 1 | 7.8           | 0.88             | 0.00        | 2.6            | 0.098     | 25.0                | 67.0                 | 0.9968  | 3.20 | 0.68      | 10.5    |
| 2 | 7.8           | 0.76             | 0.04        | 2.3            | 0.092     | 15.0                | 54.0                 | 0.9970  | 3.26 | 0.65      | 10.5    |
| 3 | 11.2          | 0.28             | 0.56        | 1.9            | 0.075     | 17.0                | 60.0                 | 0.9980  | 3.16 | 0.58      | 10.5    |
| 4 | 7.4           | 0.70             | 0.00        | 1.9            | 0.076     | 11.0                | 34.0                 | 0.9978  | 3.51 | 0.56      | 10.5    |

```
In [19]: df.info()  
df.isnull().sum()  
df.describe()
```

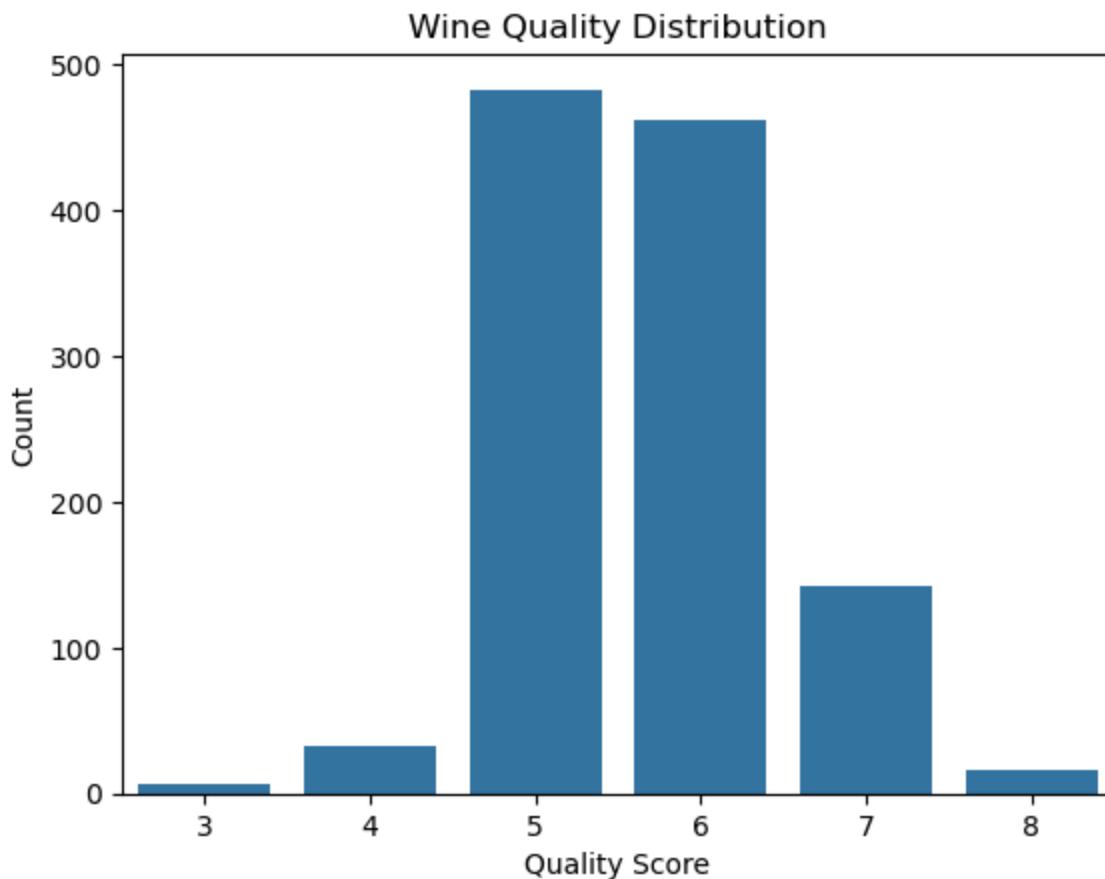
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1143 entries, 0 to 1142
Data columns (total 13 columns):
 #   Column           Non-Null Count Dtype  
 --- 
 0   fixed acidity    1143 non-null   float64 
 1   volatile acidity 1143 non-null   float64 
 2   citric acid      1143 non-null   float64 
 3   residual sugar   1143 non-null   float64 
 4   chlorides        1143 non-null   float64 
 5   free sulfur dioxide 1143 non-null   float64 
 6   total sulfur dioxide 1143 non-null   float64 
 7   density          1143 non-null   float64 
 8   pH               1143 non-null   float64 
 9   sulphates        1143 non-null   float64 
 10  alcohol          1143 non-null   float64 
 11  quality          1143 non-null   int64  
 12  Id               1143 non-null   int64  
dtypes: float64(11), int64(2)
memory usage: 116.2 KB
```

Out[19]:

|              | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides   | free sulfur dioxide | total sulfur dioxide |
|--------------|---------------|------------------|-------------|----------------|-------------|---------------------|----------------------|
| <b>count</b> | 1143.000000   | 1143.000000      | 1143.000000 | 1143.000000    | 1143.000000 | 1143.000000         | 1143.000000          |
| <b>mean</b>  | 8.311111      | 0.531339         | 0.268364    | 2.532152       | 0.086933    | 15.615486           | 45.500000            |
| <b>std</b>   | 1.747595      | 0.179633         | 0.196686    | 1.355917       | 0.047267    | 10.250486           | 32.000000            |
| <b>min</b>   | 4.600000      | 0.120000         | 0.000000    | 0.900000       | 0.012000    | 1.000000            | 6.000000             |
| <b>25%</b>   | 7.100000      | 0.392500         | 0.090000    | 1.900000       | 0.070000    | 7.000000            | 21.000000            |
| <b>50%</b>   | 7.900000      | 0.520000         | 0.250000    | 2.200000       | 0.079000    | 13.000000           | 37.000000            |
| <b>75%</b>   | 9.100000      | 0.640000         | 0.420000    | 2.600000       | 0.090000    | 21.000000           | 61.000000            |
| <b>max</b>   | 15.900000     | 1.580000         | 1.000000    | 15.500000      | 0.611000    | 68.000000           | 289.000000           |

In [20]:

```
sns.countplot(x='quality', data=df)
plt.title("Wine Quality Distribution")
plt.xlabel("Quality Score")
plt.ylabel("Count")
plt.show()
```



In [21]:

```
na_counts = df.isnull().sum()
print("Missing values per column:\n", na_counts)
```

```

non_numeric_cols = df.select_dtypes(exclude=[np.number]).columns.tolist()
print("Non-numeric columns:", non_numeric_cols)

for col in non_numeric_cols:
    df[col] = pd.to_numeric(df[col], errors='coerce')

if df.isnull().sum().sum() > 0:
    df = df.fillna(df.median(numeric_only=True))

print("Quality distribution:\n", df['quality'].value_counts().sort_index())

```

Missing values per column:

|                      |   |
|----------------------|---|
| fixed acidity        | 0 |
| volatile acidity     | 0 |
| citric acid          | 0 |
| residual sugar       | 0 |
| chlorides            | 0 |
| free sulfur dioxide  | 0 |
| total sulfur dioxide | 0 |
| density              | 0 |
| pH                   | 0 |
| sulphates            | 0 |
| alcohol              | 0 |
| quality              | 0 |
| Id                   | 0 |

dtype: int64

Non-numeric columns: []

Quality distribution:

|         |     |
|---------|-----|
| quality |     |
| 3       | 6   |
| 4       | 33  |
| 5       | 483 |
| 6       | 462 |
| 7       | 143 |
| 8       | 16  |

Name: count, dtype: int64

In [22]: `df['quality_binary'] = df['quality'].apply(lambda x: 1 if x >= 6 else 0)`

In [23]: `# Features and target`  
`X = df.drop('quality', axis=1)`  
`y = df['quality']`  
  
`# Split into train and test`  
`from sklearn.model_selection import train_test_split`  
`X_train, X_test, y_train, y_test = train_test_split(`  
 `X, y, test_size=0.2, random_state=42, stratify=y # stratify على التوزيع`  
`)`

In [24]: `from sklearn.preprocessing import StandardScaler`  
  
`scaler = StandardScaler()`  
`X_scaled = scaler.fit_transform(X)`  
  
`X_train, X_test, y_train, y_test = train_test_split(`

```
X_scaled, y, test_size=0.2, random_state=42, stratify=y  
)
```

```
In [25]:  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.linear_model import SGDClassifier  
from sklearn.svm import SVC  
from sklearn.metrics import accuracy_score, classification_report  
  
# 1) Random Forest  
rf_model = RandomForestClassifier(random_state=42)  
rf_model.fit(X_train, y_train)  
rf_pred = rf_model.predict(X_test)  
print("Random Forest Accuracy:", accuracy_score(y_test, rf_pred))  
print(classification_report(y_test, rf_pred))  
  
# 2) SGD Classifier  
sgd_model = SGDClassifier(random_state=42)  
sgd_model.fit(X_train, y_train)  
sgd_pred = sgd_model.predict(X_test)  
print("SGD Accuracy:", accuracy_score(y_test, sgd_pred))  
print(classification_report(y_test, sgd_pred))  
  
# 3) Support Vector Classifier  
svc_model = SVC(random_state=42)  
svc_model.fit(X_train, y_train)  
svc_pred = svc_model.predict(X_test)  
print("SVC Accuracy:", accuracy_score(y_test, svc_pred))  
print(classification_report(y_test, svc_pred))
```

Random Forest Accuracy: 0.8646288209606987

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 3            | 0.00      | 0.00   | 0.00     | 1       |
| 4            | 0.00      | 0.00   | 0.00     | 7       |
| 5            | 0.92      | 1.00   | 0.96     | 97      |
| 6            | 0.83      | 0.96   | 0.89     | 92      |
| 7            | 0.72      | 0.45   | 0.55     | 29      |
| 8            | 0.00      | 0.00   | 0.00     | 3       |
| accuracy     |           |        | 0.86     | 229     |
| macro avg    | 0.41      | 0.40   | 0.40     | 229     |
| weighted avg | 0.82      | 0.86   | 0.83     | 229     |

SGD Accuracy: 0.7729257641921398

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 3            | 0.00      | 0.00   | 0.00     | 1       |
| 4            | 0.00      | 0.00   | 0.00     | 7       |
| 5            | 0.92      | 0.99   | 0.96     | 97      |
| 6            | 0.84      | 0.72   | 0.77     | 92      |
| 7            | 0.35      | 0.52   | 0.42     | 29      |
| 8            | 0.00      | 0.00   | 0.00     | 3       |
| accuracy     |           |        | 0.77     | 229     |
| macro avg    | 0.35      | 0.37   | 0.36     | 229     |
| weighted avg | 0.77      | 0.77   | 0.77     | 229     |

SVC Accuracy: 0.8471615720524017

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 3            | 0.00      | 0.00   | 0.00     | 1       |
| 4            | 0.00      | 0.00   | 0.00     | 7       |
| 5            | 0.91      | 1.00   | 0.95     | 97      |
| 6            | 0.80      | 0.95   | 0.87     | 92      |
| 7            | 0.77      | 0.34   | 0.48     | 29      |
| 8            | 0.00      | 0.00   | 0.00     | 3       |
| accuracy     |           |        | 0.85     | 229     |
| macro avg    | 0.41      | 0.38   | 0.38     | 229     |
| weighted avg | 0.80      | 0.85   | 0.81     | 229     |

```
D:\Anaconda_Set_Up\Lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
D:\Anaconda_Set_Up\Lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
D:\Anaconda_Set_Up\Lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
D:\Anaconda_Set_Up\Lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
D:\Anaconda_Set_Up\Lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
D:\Anaconda_Set_Up\Lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
D:\Anaconda_Set_Up\Lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
D:\Anaconda_Set_Up\Lib\site-packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    warn prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
In [26]: from sklearn.ensemble import GradientBoostingClassifier
```

```
gb_model = GradientBoostingClassifier(random_state=42)
gb_model.fit(X_train, y_train)
gb_pred = gb_model.predict(X_test)

print("Gradient Boosting Accuracy:", accuracy_score(y_test, gb_pred))
print(classification_report(y_test, gb_pred))
```

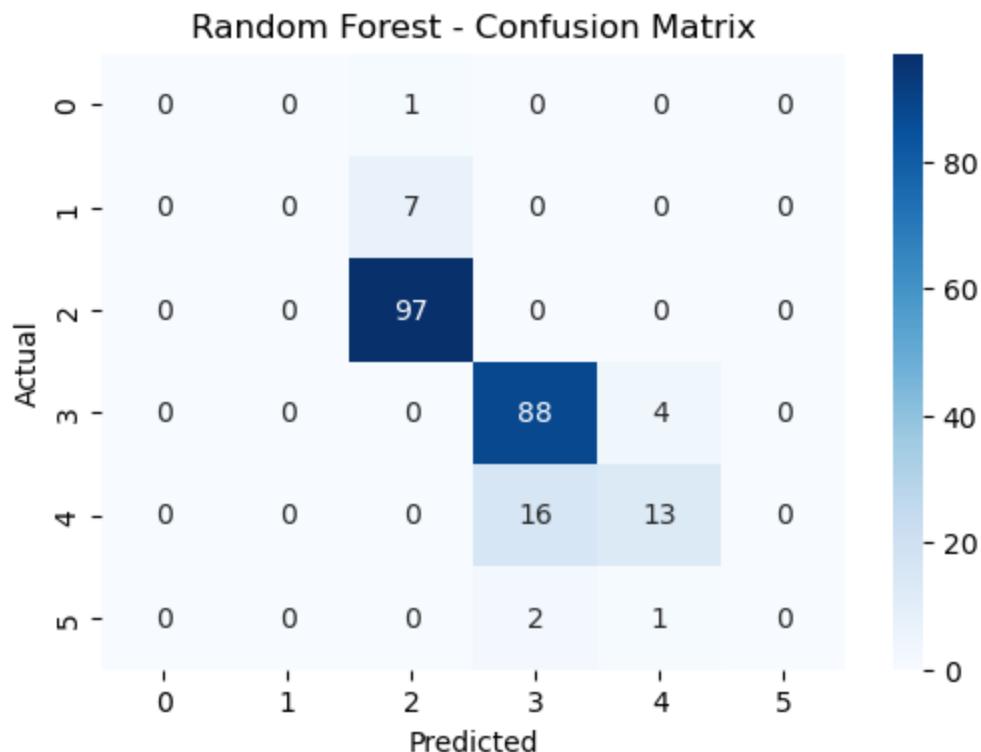
Gradient Boosting Accuracy: 0.8646288209606987

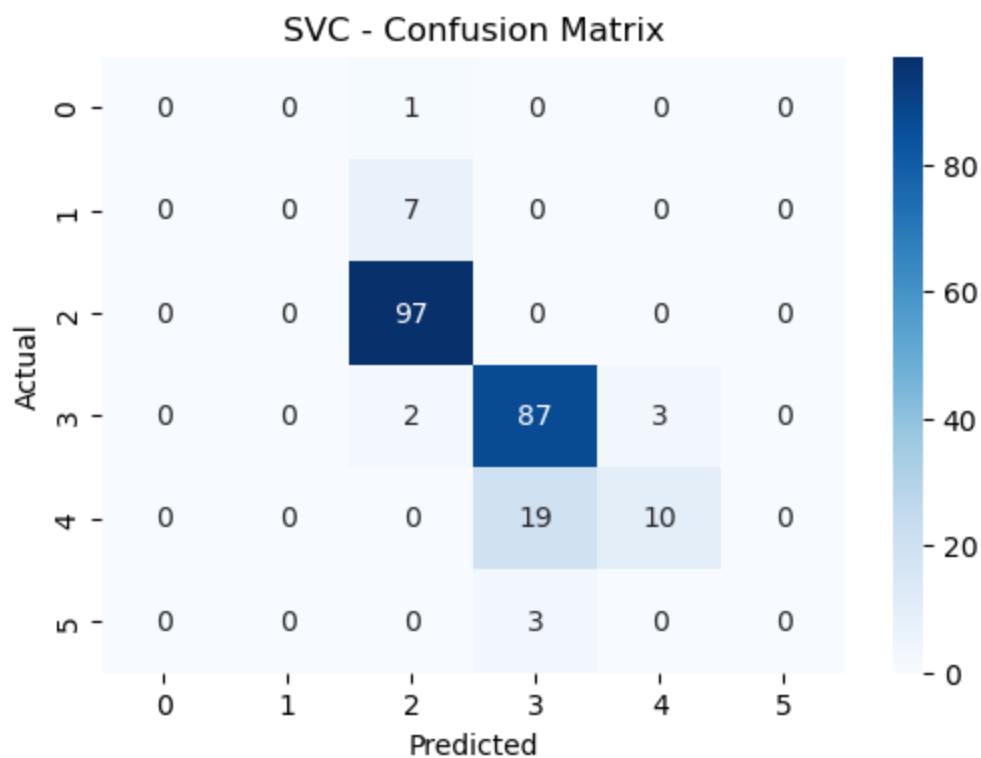
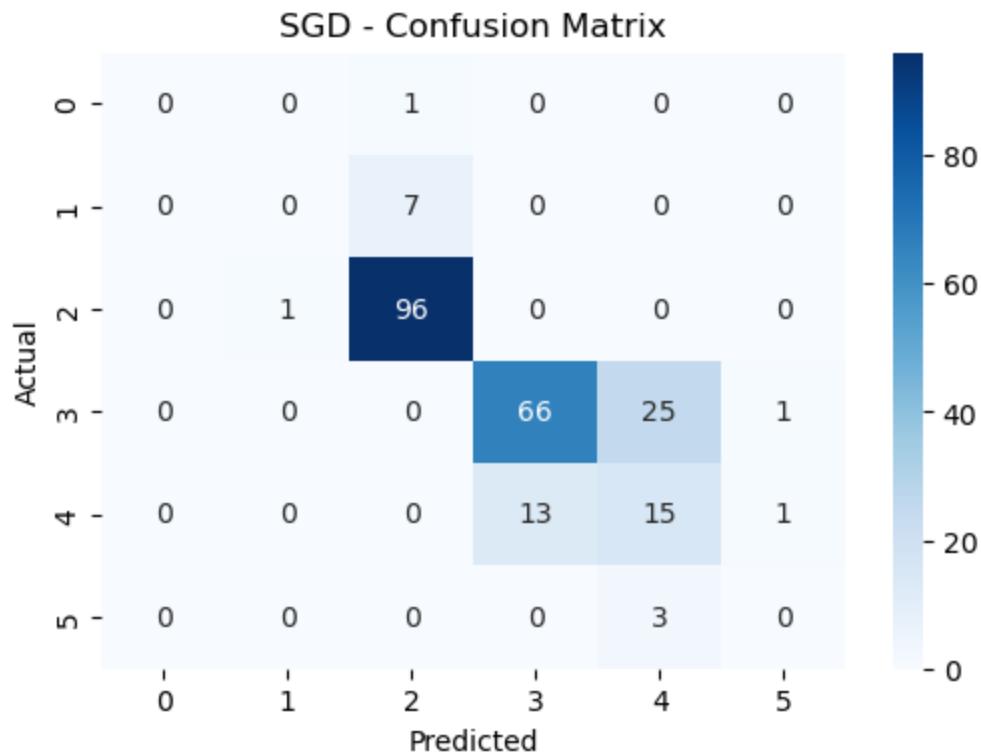
|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 3            | 0.00      | 0.00   | 0.00     | 1       |
| 4            | 0.00      | 0.00   | 0.00     | 7       |
| 5            | 0.93      | 0.99   | 0.96     | 97      |
| 6            | 0.85      | 0.92   | 0.89     | 92      |
| 7            | 0.74      | 0.59   | 0.65     | 29      |
| 8            | 0.00      | 0.00   | 0.00     | 3       |
| accuracy     |           |        | 0.86     | 229     |
| macro avg    | 0.42      | 0.42   | 0.42     | 229     |
| weighted avg | 0.83      | 0.86   | 0.85     | 229     |

```
In [27]: import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix
```

```
def plot_confusion(y_true, y_pred, title):
    cm = confusion_matrix(y_true, y_pred)
    plt.figure(figsize=(6,4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
    plt.title(title)
    plt.xlabel('Predicted')
    plt.ylabel('Actual')
    plt.show()

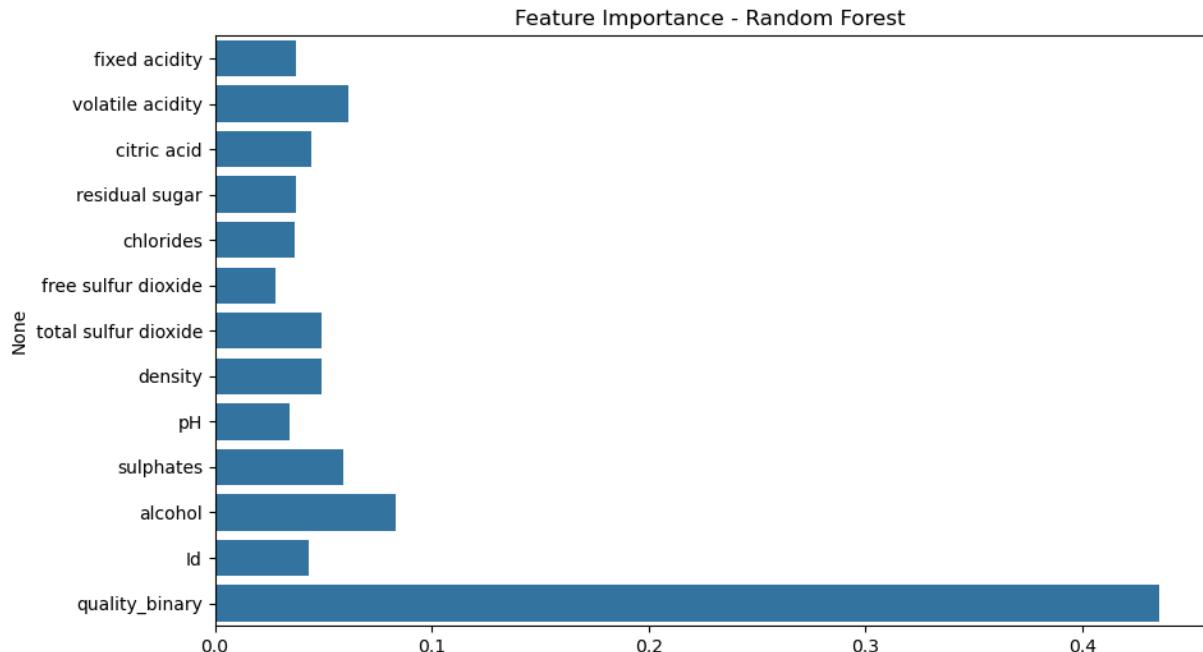
plot_confusion(y_test, rf_pred, "Random Forest - Confusion Matrix")
plot_confusion(y_test, sgd_pred, "SGD - Confusion Matrix")
plot_confusion(y_test, svc_pred, "SVC - Confusion Matrix")
```





```
In [28]: importances = rf_model.feature_importances_
features = X.columns

plt.figure(figsize=(10,6))
sns.barplot(x=importances, y=features)
plt.title("Feature Importance - Random Forest")
plt.show()
```



```
In [29]: results = {
    "RandomForest": accuracy_score(y_test, rf_pred),
    "SGD": accuracy_score(y_test, sgd_pred),
    "SVC": accuracy_score(y_test, svc_pred)
}
print("Model accuracies:", results)
best_model = max(results, key=results.get)
print("Best model:", best_model)
```

Model accuracies: {'RandomForest': 0.8646288209606987, 'SGD': 0.7729257641921398, 'SVC': 0.8471615720524017}  
Best model: RandomForest